

# StackGuard: Automatic Adaptive Detection and Prevention of Buffer-Overflow Attacks

Will Henderson, Josh Dall'Acqua

COMP 310, McGill University

February 16, 2024

# The Morris Worm

In 1988, Cornell graduate student Robert Morris released the first mainstream **computer worm** to propagate through the Internet, infecting an estimated 10% of UNIX systems connected at the time.

Though intending to be merely an experiment of how far the worm could spread, it effectively took down a good handful of computers.

# Morris's Mistake

Morris knew that when he checked if a computer was already infected with the worm, some people may send false positives back. To get around this and keep the worm propagating, he tacked on a **14%** chance of infecting a computer that sent a positive signal.

# Morris's Mistake

Morris knew that when he checked if a computer was already infected with the worm, some people may send false positives back. To get around this and keep the worm propagating, he tacked on a **14%** chance of infecting a computer that sent a positive signal.

Unfortunately, this meant some computers got infected many, many times...

# The Morris Worm

This was a big deal. It caused monetary damages, lost people's trust in the Internet, and led to a conviction for Morris.

# The Morris Worm

This was a big deal. It caused monetary damages, lost people's trust in the Internet, and led to a conviction for Morris.

...How exactly did he do it, though?

# The Morris Worm

One of the main exploits he used was **buffer overflows**.

This presentation introduces the buffer overflow exploit and *StackGuard*, a 1998 paper that details a way to patch it.

# Table of Contents

- 1 Introduction
- 2 Background
- 3 Methodology
  - Stack Guard
  - Mem Guard
- 4 Results
- 5 Presenters' Thoughts
- 6 Conclusion



# Table of Contents

- 1 Introduction
- 2 Background
- 3 Methodology
  - Stack Guard
  - Mem Guard
- 4 Results
- 5 Presenters' Thoughts
- 6 Conclusion

# Problem

- Buffer overflow vulnerabilities are common in legacy code
- It is easy to patch a vulnerability but expensive to find them
- What is a low cost, general way to prevent these vulnerabilities?

- Stack Guard aims to provide a low cost, effective solution to prevent buffer overflow attacks
  - Accomplished through a compiler patch
  - Does not modify source code
  - Focuses on preventing changes to a program's return address.

# Ramifications

- Deescalate urgency in patching buffer overflow vulnerabilities
- Immediately increase confidence in security of existing code

# Difficulties

- Adding detection of buffer overflow attacks is expensive
- Implementation must be general and work with existing code

# Table of Contents

1 Introduction

2 Background

3 Methodology

- Stack Guard
- Mem Guard

4 Results

5 Presenters' Thoughts


6 Conclusion

# Connection to Course Content

- Pages and Virtual Pages
- Layout of the stack in memory

# What is a Buffer Overflow Attack?

- C and C++ **do not** provide standard array bounds checking
- Data can be written into memory after the end of the array (potentially outside of the program's memory block!)
- **Exploitation opportunity**: code can be injected into memory after the array



buffer array boundary



# Stack Smashing Attacks

- A very common overflow attack is the ingeniously simple **stack smashing attack**, performed by providing an input string that works in two parts:
  - Malicious binary code that is executable on the machine being attacked, i.e., **attack code**, is injected onto the stack beyond the buffer, and
  - The **return address** of the function, also contained in a stack frame beyond the buffer, is modified to then jump to the attack code.
- Most often, this type of attack is used against privileged daemons to create a new root shell, providing the attacker with root privileges.

# Table of Contents

1 Introduction

2 Background

3 Methodology

- Stack Guard
- Mem Guard

4 Results

5 Presenters' Thoughts

6 Conclusion

# How Does Stack Guard Work?

- Prevent injected code from changing the function's return address
- Stops running program (potentially on a privileged daemon) from entering injected attack code.
- Two approaches with different performance trade-offs
  - Canary word
  - Memory Guard
- Both approaches modify gcc to check the integrity of the function's return address before returning

# Canary Stack Guard

- Place a "canary" word next to the return address on the stack




Figure: Canary word location in stack, from [1].

- If the attack code attempts to change the return address via overflow, the overflow will overwrite the canary

# Canary Stack Guard

- Before returning, the function epilogue checks that the canary word is intact
- Randomizing the canary
  - A constant canary is easy to guess depending on how much access the attacker has
  - It is then easy to bypass Stack Guard
  - This can be avoided by randomly choosing a canary word

# Virtual Memory Mem Guard

- Protect function return address when the function is called and un-protect it when the function returns
- The page in virtual memory where the return address is stored must be marked read-only
- Mem Guard API will emulate writes to non protected words on the page while keeping the return address protected
- Comes with a larger performance penalty than the canary word approach
  - Emulating Writes to non protected words is very expensive

# Pentium Debug Register Mem Guard

- Pentium processors have 4 debug registers that can watch for reads/writes to the addresses stored in each register
- Can be configured to generate an exception when a write is made to one of the stored addresses
- By storing most recently used return addresses in the debug registers, writes to those registers can be prevented
- This strategy is more performant than the previously discussed virtual memory approach

# Adaptive Defense Strategy

- Mem Guard implementation is significantly more expensive than using a canary word
  - Discussed in results section
- Adaptive strategy will use "canary mode" until an attack is detected at which point it will switch to Mem Guard
- Performance is only significantly degraded during attack but security remains high



# Table of Contents

- 1 Introduction
- 2 Background
- 3 Methodology
  - Stack Guard
  - Mem Guard
- 4 Results
- 5 Presenters' Thoughts
- 6 Conclusion

# Attack Prevention Experiments

- Experimental penetrative tests were run on programs with and without StackGuard to gauge effectiveness
- Methodology involved simulating generic attacks on roots
- General success found with both Canary and MemGuard versions

- Both highly effective
- Even for unprecedented, new attack styles
- Some effectiveness shown on similar attacks (namely Perl), but vulnerabilities when attacking things other than the function return address

# Results

Vulnerable Program	Result Without StackGuard	Result With Canary StackGuard	Result With MemGuard StackGuard
dip 3.3.7n	root shell	program halts	program halts
elm 2.4 PL25	root shell	program halts	program halts
Perl 5.003	root shell	program halts irregularly	root shell
Samba	root shell	program halts	program halts
SuperProbe	root shell	program halts irregularly	program halts
umount 2.5k/libc 5.3.12	root shell	program halts	program halts
wwwcount v2.3	httpd shell	program halts	program halts
zgv 2.7	root shell	program halts	program halts

Figure: Results from experiments on various programs, from [1]

- Of course, implementing added protection requires some overhead
- Experimental results were taken to find the ratio (%) of additional overhead relative to the non-StackGuard version of the program

# Canary Overhead

- Pushing the canary onto the stack
- Checking the canary word is intact

# Canary Overhead

Increment Method	Standard Run-Time	Canary Run-Time	% Overhead
<code>i++</code>	15.1	15.1	NA
<code>void inc()</code>	35.1	60.2	125%
<code>void inc(int *)</code>	47.7	70.2	69%
<code>int inc(int)</code>	40.1	60.2	80%

Figure: Canary overhead testing results

# MemGuard Overhead

Significantly worse than canary

Increment Method	Standard Run-Time	MemGuard Register Run-Time	% Overhead	MemGuard VM Run-Time	% Overhead
<code>i++</code>	15.1	15.1	NA	NA	NA
<code>void inc()</code>	35.1	1808	8800%	34,900	174,300%
<code>void inc(int *)</code>	47.7	1820	5400%	40,420	123,800%
<code>int inc(int)</code>	40.1	1815	7000%	41,610	166,200%

Figure: MemGuard overhead testing results



# Table of Contents

- 1 Introduction
- 2 Background
- 3 Methodology
  - Stack Guard
  - Mem Guard
- 4 Results
- 5 Presenters' Thoughts
- 6 Conclusion

# Overheads and Compromises

- It goes without saying that the MemGuard overhead is ridiculously bad
- In comparison, canary looks great, but alone, it is still almost twice as slow in cases
- As a result, we should strive to only use this version of StackGuard when absolutely needed

# Overheads and Compromises

- Luckily, more recent developments in the area of buffer overflow prevention have been working on reconciling this

# Table of Contents

- 1 Introduction
- 2 Background
- 3 Methodology
  - Stack Guard
  - Mem Guard
- 4 Results
- 5 Presenters' Thoughts
- 6 Conclusion

# Timeline of Protection in the GCC

1998: *StackGuard* is published

2001: IBM creates *ProPolice*, inspired by StackGuard

2005: A derivative of ProPolice is included in GCC 4.1

- `-fstack-protector` and `-fstack-protector-all`

2012: GCC 4.9 adds `-fstack-protector-strong`

These days, many Linux distributions (incl. Arch, Fedora, and Ubuntu) packages are compiled with this protection by default

# StackGuard: a Pioneer?

- Indeed, StackGuard was the first of its kind
- Never actually implemented in GCC, but successors were
- This was a great first step, as it set the stage for buffer overflow attack prevention

- [1] Crispian Cowan et al. “StackGuard: Automatic Adaptive Detection and Prevention of Buffer-Overflow Attacks”. In: *USENIX security symposium 98* (1998), pp. 63–78.